

K-Means

```
In [70]: pip install -U numpy
```

Requirement already satisfied: numpy in /Users/haoranzhang/opt/anaconda3/lib/python3.8/site-packages (1.24.2)
Note: you may need to restart the kernel to use updated packages.

```
In [71]: import sklearn
```

```
In [72]: from sklearn.cluster import KMeans
```

```
In [73]: import sklearn.cluster as cluster
```

```
In [74]: import matplotlib.pyplot as plt  
import math
```

```
In [75]: import pandas as pd
```

```
In [76]: import numpy as np
```

Q1 (a)(b)

```
In [77]: data = np.loadtxt("/Users/haoranzhang/CS6220/assignment-4-hrcchung/data/f15
```

```
In [78]: data
```

```
Out[78]: array([[ -11.96999577,  -8.03962819],  
                [-26.96141582,  -6.96210943],  
                [-12.91584891,  -1.37894148],  
                ...,  
                [ 12.92550271,  -0.10319682],  
                [-34.33677838,   4.4057531 ],  
                [-26.81273855,  -1.63995599]])
```

```
In [79]: data.shape
```

```
Out[79]: (5000, 2)
```

```
In [80]: #initialization vectors are:  
initialization=np.array([[10,10],[-10,-10],[2,2],[3,3],[-3,-3]])
```

```
In [81]: type(initialization)
```

```
Out[81]: numpy.ndarray
```

```
In [82]: initialization.shape
```

```
Out[82]: (5, 2)
```

```
In [83]: kmeans=cluster.KMeans(n_clusters=5,init=initialization,max_iter=100)
```

```
In [84]: label=kmeans.fit_predict(data)
```

```
/Users/haoranzhang/opt/anaconda3/lib/python3.8/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
/Users/haoranzhang/opt/anaconda3/lib/python3.8/site-packages/sklearn/cluster/_kmeans.py:1362: RuntimeWarning: Explicit initial center position passed: performing only one init in KMeans instead of n_init=10.
  super()._check_params_vs_input(X, default_n_init=10)
```


```
In [85]: label
```

```
Out[85]: array([4, 1, 4, ..., 3, 1, 1], dtype=int32)
```

```
In [86]: centroids = kmeans.cluster_centers_
centroids
```

```
Out[86]: array([[ 27.3664243 , -1.09580993],
                [-31.94805334,  2.59901959],
                [ -1.04061308, -2.28300242],
                [ 11.68619774, -1.68891573],
                [-14.50973012, -0.54232205]])
```

```
In [87]: u_labels=np.unique(label) #allocat label to different clusters
for i in u_labels:
    plt.scatter(data[label==i,0],data[label==i,1],label=i)
plt.scatter(centroids[:,0] , centroids[:,1] , s = 80, color = 'k')
plt.legend()
plt.show()
```



Q1(c)


Based on the cluster figure above, the cluster result does not look good.

I am suspecting the initialization vectors and centroids are incorrect. They are all on the line of $y=x$, which disturbed the accurate clustering process.

Based on the picture, a better initialization would be $y=c$, where c is the constant, since apparently points can be clustered as blocks layered parallel to the x axis.

```
In [88]: cost = []
K = range(1,10)
for k in K:
    kmeanModel = cluster.KMeans(n_clusters=k)
    kmeanModel.fit(data)
    cost.append(kmeanModel.inertia_) #sum of distances to cluster center

plt.figure(figsize=(16,8))
plt.plot(K, cost, 'bx-')
plt.xlabel('K number of clusters')
plt.ylabel('Cost')
plt.title('Elbow method')
plt.show()
```

k=3 is a relatively good choice from the perspective of cost, according to Elbow method.

Q1 a) b) write k-means by hand

```
In [89]: df = pd.DataFrame(data, columns = ['x', 'y'])
df
```

Out[89]:

	x	y
0	-11.969996	-8.039628
1	-26.961416	-6.962109
2	-12.915849	-1.378941
3	22.476144	2.066612
4	-13.146631	4.835322
...
4995	6.852668	1.549076
4996	-2.340729	-7.343469
4997	12.925503	-0.103197
4998	-34.336778	4.405753
4999	-26.812739	-1.639956

5000 rows × 2 columns

```

In [90]: def euclideanDistance(p1,p2):
            return math.sqrt(((p1[0]-p2[0])**2)+((p1[1]-p2[1])**2) )

def find_closest_centroid(point,centroids):
    # for i in range(len(centroids)):
    #     clusters[i]=list()
    minV=float("inf")
    label=-1
    for i, center in enumerate(centroids):
        if euclideanDistance(center,point)<minV:
            label=i
            minV=euclideanDistance(center,point)
    return label

#compute new centroid after assign
def computeNewCentroid(X):
    new=list()
    tmp=X.groupby(by='label').mean()
    for row in range(len(tmp)):
        new.append([tmp.at[row,'x'],tmp.at[row,'y']])
    return new

#check if centroids change
def centroidsChange(old,new):
    if euclideanDistance(old,new)>0.000000000001:
        return True
    else:
        return False

def fit(X,initialization,max_iter,n_cluster):
    centroids=initialization
    # print("initial centroids =",centroids)

    X['label']=X.apply(lambda row: find_closest_centroid((row[0],row[1]),centroids),axis=1)
    for i in range(max_iter):
        new_centroids=computeNewCentroid(X)
        # print("new centroids is",new_centroids)
        for i in range(len(centroids)):
            if not centroidsChange(centroids[i],new_centroids[i]):
                break
            else:
                centroids=new_centroids
                X['label']=X.apply(lambda row: find_closest_centroid((row[0],row[1]),centroids),axis=1)
    return centroids

initialization=[[10,10],[-10,-10],[2,2],[3,3],[-3,-3]]
# new_centroids=findNewCentroids(centroids)
clusters=fit(df,initialization,100,5)
clusters=np.asarray(clusters)
clusters

```

```
Out[90]: array([[ 28.09146831, -1.0241717 ],
                [-30.93833893,  2.43436785],
                [  0.52511293, -2.3158082 ],
                [ 12.86956124, -1.68099097],
                [-12.88359021, -0.80312447]])
```


```
In [91]: df
```

```
Out[91]:
```

	x	y	label
0	-11.969996	-8.039628	4
1	-26.961416	-6.962109	1
2	-12.915849	-1.378941	4
3	22.476144	2.066612	0
4	-13.146631	4.835322	4
...
4995	6.852668	1.549076	3
4996	-2.340729	-7.343469	2
4997	12.925503	-0.103197	3
4998	-34.336778	4.405753	1
4999	-26.812739	-1.639956	1

5000 rows × 3 columns

```
In [92]: u_labels=np.unique(df['label']) #allocat label to different clusters
for i in u_labels:
    plt.scatter(data[label==i,0],data[label==i,1],label=i)
plt.scatter(clusters[:,0] , clusters[:,1] , s = 80, color = 'k')
plt.legend()
plt.show()
```



Q2

Q2 a)

```
In [93]: #compute P-1 first
P_value=[[10,0.5],[-10,0.25]]
P=np.matrix(P_value)
R=(P.T @ P).I
```

```
In [94]: R
```

```
Out[94]: matrix([[ 0.00555556, -0.04444444],
                 [-0.04444444,  3.55555556]])
```

```

In [95]: def mahalaDistance(x,y):
            x=np.asarray(x)
            y=np.asarray(y)
            return (x-y).T @ R @ (x-y)

def find_closest_centroid(point,centroids):
    # for i in range(len(centroids)):
    #     clusters[i]=list()
    minV=float("inf")
    label=-1
    for i, center in enumerate(centroids):
        if mahalaDistance(center,point)<minV:
            label=i
            minV=mahalaDistance(center,point)
    return label

#compute new centroid after assign
def computeNewCentroid(X):
    new=list()
    tmp=X.groupby(by='label').mean()
    for row in range(len(tmp)):
        new.append([tmp.at[row,'x'],tmp.at[row,'y']])
    return new

#check if centroids change
def centroidsChange(old,new):
    if mahalaDistance(old,new)>0.000000000001:
        return True
    else:
        return False

def fit(X,initialization,max_iter,n_cluster):
    centroids=initialization
    # print("initial centroids =",centroids)

    X['label']=X.apply(lambda row: find_closest_centroid((row[0],row[1]),centroids),axis=1)
    for i in range(max_iter):
        new_centroids=computeNewCentroid(X)
        # print("new centroids is",new_centroids)
        for i in range(len(centroids)):
            if not centroidsChange(centroids[i],new_centroids[i]):
                break
            else:
                centroids=new_centroids
                X['label']=X.apply(lambda row: find_closest_centroid((row[0],row[1]),centroids),axis=1)
    return centroids

```

```

In [96]: df= pd.DataFrame(data, columns = ['x','y'])

```

```
In [97]: df
```

```
Out[97]:
```

	x	y
0	-11.969996	-8.039628
1	-26.961416	-6.962109
2	-12.915849	-1.378941
3	22.476144	2.066612
4	-13.146631	4.835322
...
4995	6.852668	1.549076
4996	-2.340729	-7.343469
4997	12.925503	-0.103197
4998	-34.336778	4.405753
4999	-26.812739	-1.639956

5000 rows × 2 columns

```
In [98]: initialization=[[10,10],[-10,-10],[2,2],[3,3],[-3,-3]]
#           new_centroids=findNewCentroids(centroids)
clusters=fit(df,initialization,100,5)
clusters=np.asarray(clusters)
clusters
```

```
Out[98]: array([[ -20.40310599,   5.00655283],
 [  -3.24764066,  -7.01766444],
 [  -1.00542627,  -0.97629543],
 [   9.96452001,   2.01588007],
 [   8.72391858,  -4.00459056]])
```


```
In [99]: df
```

```
Out[99]:
```

	x	y	label
0	-11.969996	-8.039628	1
1	-26.961416	-6.962109	1
2	-12.915849	-1.378941	2
3	22.476144	2.066612	3
4	-13.146631	4.835322	0
...
4995	6.852668	1.549076	3
4996	-2.340729	-7.343469	1
4997	12.925503	-0.103197	2
4998	-34.336778	4.405753	0
4999	-26.812739	-1.639956	2

5000 rows × 3 columns

```
In [100]: u_labels=np.unique(df['label']) #allocat label to different clusters
for i in u_labels:
    row=np.where(df.label==i)
    plt.scatter(df.iloc[row]['x'],df.iloc[row]['y'],label=i)
plt.scatter(clusters[:,0] , clusters[:,1] , s = 80, color = 'k')
plt.legend()
plt.show()
```



Points are clustered into a right way. Using Mahalanobis Distance allows a different way to calculate similarity between points.

Q2 b)

```
In [101]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
import sklearn
```

```
In [102]: data_q2=df
data_q2.head()
```

```
Out[102]:
```

	x	y	label
0	-11.969996	-8.039628	1
1	-26.961416	-6.962109	1
2	-12.915849	-1.378941	2
3	22.476144	2.066612	3
4	-13.146631	4.835322	0

```
In [103]: target=data_q2['label']
```

```
In [104]: data=data_q2.loc[:,['x','y']]
```

```
In [105]: data
```

```
Out[105]:
```

	x	y
0	-11.969996	-8.039628
1	-26.961416	-6.962109
2	-12.915849	-1.378941
3	22.476144	2.066612
4	-13.146631	4.835322
...
4995	6.852668	1.549076
4996	-2.340729	-7.343469
4997	12.925503	-0.103197
4998	-34.336778	4.405753
4999	-26.812739	-1.639956

5000 rows × 2 columns

```
In [106]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
print(scaler.fit(data))
```

```
StandardScaler()
```

```
In [107]: print(scaler.mean_)
```

```
[-1.22012285 -1.00360882]
```

```
In [108]: data_z_score=scaler.transform(data)
```

```
In [109]: print(data_z_score)
```

```
[[-0.59957277 -1.64159106]
 [-1.43571729 -1.39019249]
 [-0.65232761 -0.08756979]
 ...
 [ 0.78897044  0.21007735]
 [-1.84707718  1.26207159]
 [-1.42742483 -0.14846773]]
```

PCA

```
In [110]: from sklearn.decomposition import PCA
# Two components of PCA
pca = PCA(2)

# Fit on data
pca.fit(data_z_score)

# Access values and vectors
print(pca.components_)
print(pca.explained_variance_)

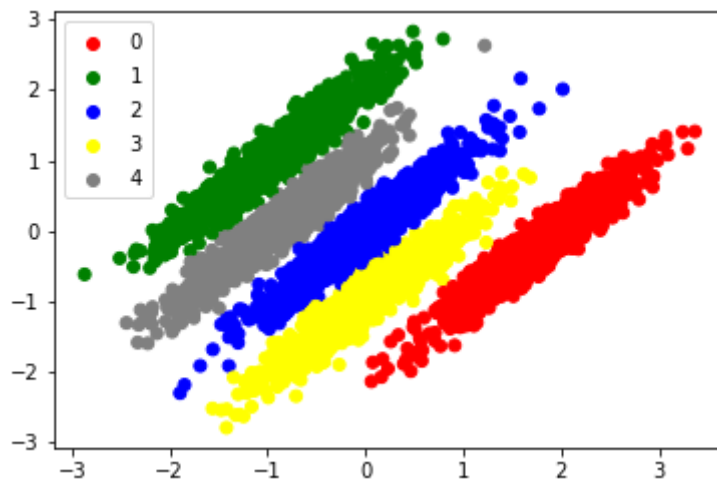
# transform data
B = pca.transform(data_z_score)
print(B)
```

```
[[-0.70710678  0.70710678]
 [-0.70710678 -0.70710678]]
[1.22552735  0.77487273]
[[-0.7368182   1.58474214]
 [ 0.03219089  1.99821997]
 [ 0.39934408  0.52318647]
 ...
 [-0.40933923 -0.70643346]
 [ 2.19850018  0.41366142]
 [ 0.90435924  1.11432431]]
```

```
In [111]: B.shape
```

```
Out[111]: (5000, 2)
```

```
In [112]: import matplotlib.pyplot as plt
import matplotlib
colors=['red','green','blue','yellow','grey']
scatter=plt.scatter(B[:, 0], B[:, 1],c=target,\
                    cmap=matplotlib.colors.ListedColormap(colors),marker='o')
plt.legend(*scatter.legend_elements())
plt.show()
```



if not Z-score scalar

```
In [113]: data
```

```
Out[113]:
```

	x	y
0	-11.969996	-8.039628
1	-26.961416	-6.962109
2	-12.915849	-1.378941
3	22.476144	2.066612
4	-13.146631	4.835322
...
4995	6.852668	1.549076
4996	-2.340729	-7.343469
4997	12.925503	-0.103197
4998	-34.336778	4.405753
4999	-26.812739	-1.639956

5000 rows × 2 columns

```
In [114]: from sklearn.decomposition import PCA
# Two components of PCA
pca = PCA(2)

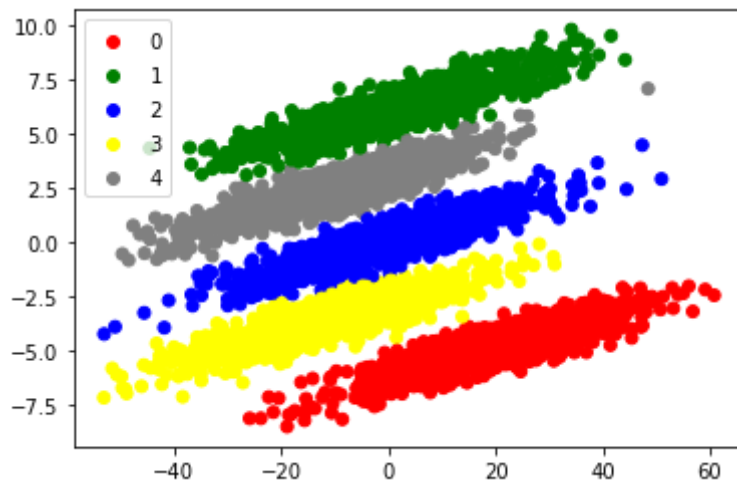
# Fit on data
pca.fit(data)

# Access values and vectors
print(pca.components_)
print(pca.explained_variance_)

# transform data
B = pca.transform(data)
print(B)
```

```
[[-0.99838317  0.05684225]
 [-0.05684225 -0.99838317]]
[322.50713273  17.38845582]
[[ 10.33254906   7.6356903 ]
 [ 25.36097916   7.41205973]
 [ 11.65548133   1.03953719]
 ...
 [-14.07157308  -1.70302537]
 [ 33.3705919   -3.51819072]
 [ 25.51506544   2.09006015]]
```

```
In [115]: import matplotlib.pyplot as plt
import matplotlib
colors=['red','green','blue','yellow','grey']
scatter=plt.scatter(B[:, 0], B[:, 1],c=target,\
                    cmap=matplotlib.colors.ListedColormap(colors),marker='o')
plt.legend(*scatter.legend_elements())
plt.show()
```



Q2 c)

```
In [116]: data_q2
```

```
Out[116]:
```

	x	y	label
0	-11.969996	-8.039628	1
1	-26.961416	-6.962109	1
2	-12.915849	-1.378941	2
3	22.476144	2.066612	3
4	-13.146631	4.835322	0
...
4995	6.852668	1.549076	3
4996	-2.340729	-7.343469	1
4997	12.925503	-0.103197	2
4998	-34.336778	4.405753	0
4999	-26.812739	-1.639956	2

5000 rows × 3 columns

```
In [117]: cluster_0=data_q2[data_q2['label'] == 0]
cluster_1=data_q2[data_q2['label'] == 1]
cluster_2=data_q2[data_q2['label'] == 2]
cluster_3=data_q2[data_q2['label'] == 3]
cluster_4=data_q2[data_q2['label'] == 4]
```

```
In [118]: cluster_0
```

```
Out[118]:
```

	x	y	label
4	-13.146631	4.835322	0
12	-38.508282	4.112202	0
19	6.843345	4.550918	0
22	-36.768240	4.173526	0
25	-16.883207	4.664295	0
...
4985	-41.441343	5.164262	0
4987	-39.028493	5.153460	0
4991	-14.425117	5.286585	0
4994	-11.915634	5.320493	0
4998	-34.336778	4.405753	0

1000 rows × 3 columns

```
In [119]: cluster_0=cluster_0.iloc[:,2]
cluster_1=cluster_1.iloc[:,2]
cluster_2=cluster_2.iloc[:,2]
cluster_3=cluster_3.iloc[:,2]
cluster_4=cluster_4.iloc[:,2]
```

```
In [120]: def pca_onto_cluster(x):
#     scaler = StandardScaler()
#     scaler.fit(data)
#     data_z_score=scaler.transform(data)
# Two components of PCA
pca = PCA(2)

# Fit on data
pca.fit(x)

# Access values and vectors
# transform data
B = pca.transform(x)
return B
```

```
In [121]: res0=pca_onto_cluster(cluster_0)
res0
```

```
Out[121]: array([[ 7.25405752,  0.25378307],
 [-18.11417956,  0.68829451],
 [ 27.23950349,  0.76561297],
 ...,
 [ 5.98078814, -0.21199726],
 [ 8.49049436, -0.21735037],
 [-13.9396063 ,  0.44222502]])
```

```
In [122]: res1=pca_onto_cluster(cluster_1)
res1
```

```
Out[122]: array([[ -8.73419248,  0.91529791],
 [-23.71132555, -0.34533961],
 [ 1.94266087, -0.75298475],
 ...,
 [ 6.37240317,  0.70787872],
 [-15.07539706, -0.44106933],
 [ 0.90286284,  0.33686259]])
```

```
In [123]: res2=pca_onto_cluster(cluster_2)
res2
```

```
Out[123]: array([[ -11.81898073, -0.22222287],
 [ -7.96831058, -0.07484901],
 [-21.787302 ,  0.26996112],
 ...,
 [-15.0302669 , -0.59978662],
 [ 14.03719841,  0.70566602],
 [-25.71812447, -0.29620696]])
```

```
In [124]: res3=pca_onto_cluster(cluster_3)
res3
```

```
Out[124]: array([[ 12.50154109, -0.080114 ],
 [ 26.88083768,  0.15701845],
 [-12.34056372, -0.61900941],
 ...,
 [ 28.29556699, -0.14442832],
 [ -8.76609494,  0.96638809],
 [ -3.12687729, -0.41684377]])
```

```
In [125]: res4=pca_onto_cluster(cluster_4)
res4
```

```
Out[125]: array([[25.99078891, -0.97033838],
 [58.70744189, -0.41801869],
 [ 8.42336008, -0.14735909],
 ...,
 [-1.4299455 , -0.83051362],
 [24.30622751, -0.48802078],
 [ 0.58512901, -0.68303187]])
```

```
In [126]: res=list()  
res.append(res0)  
res.append(res1)  
res.append(res2)  
res.append(res3)  
res.append(res4)
```

```
In [127]: res
```

```
Out[127]: [array([[ 7.25405752,  0.25378307],  
                [-18.11417956,  0.68829451],  
                [ 27.23950349,  0.76561297],  
                ...,  
                [ 5.98078814, -0.21199726],  
                [ 8.49049436, -0.21735037],  
                [-13.9396063 ,  0.44222502]])],  
array([[ -8.73419248,  0.91529791],  
        [-23.71132555, -0.34533961],  
        [ 1.94266087, -0.75298475],  
        ...,  
        [ 6.37240317,  0.70787872],  
        [-15.07539706, -0.44106933],  
        [ 0.90286284,  0.33686259]])],  
array([[ -11.81898073, -0.22222287],  
        [ -7.96831058, -0.07484901],  
        [-21.787302  ,  0.26996112],  
        ...,  
        [-15.0302669 , -0.59978662],  
        [ 14.03719841,  0.70566602],  
        [-25.71812447, -0.29620696]])],  
array([[ 12.50154109, -0.080114  ],  
        [ 26.88083768,  0.15701845],  
        [-12.34056372, -0.61900941],  
        ...,  
        [ 28.29556699, -0.14442832],  
        [ -8.76609494,  0.96638809],  
        [ -3.12687729, -0.41684377]])],  
array([[25.99078891, -0.97033838],  
        [58.70744189, -0.41801869],  
        [ 8.42336008, -0.14735909],  
        ...,  
        [-1.4299455 , -0.83051362],  
        [24.30622751, -0.48802078],  
        [ 0.58512901, -0.68303187]])]
```


```
In [128]: res[0]
```

```
Out[128]: array([[ 7.25405752,  0.25378307],  
                [-18.11417956,  0.68829451],  
                [ 27.23950349,  0.76561297],  
                ...,  
                [ 5.98078814, -0.21199726],  
                [ 8.49049436, -0.21735037],  
                [-13.9396063 ,  0.44222502]])
```

```
In [129]: res[1]
```

```
Out[129]: array([[ -8.73419248,  0.91529791],
                 [-23.71132555, -0.34533961],
                 [  1.94266087, -0.75298475],
                 ...,
                 [  6.37240317,  0.70787872],
                 [-15.07539706, -0.44106933],
                 [  0.90286284,  0.33686259]])
```

```
In [132]: import matplotlib.pyplot as plt
import matplotlib
colors=['red','green','blue','yellow','grey']
scatter=plt.scatter(B[:, 0], B[:, 1],c=target,\
                    cmap=matplotlib.colors.ListedColormap(colors),marker='o')
plt.legend(*scatter.legend_elements())
plt.show()
```



The result of 2b and 2c are the same.

Q2 d)

```
In [133]: vector,mat=np.linalg.eig(P.T @P )
```

```
In [134]: vector
```

```
Out[134]: array([200.031294,  0.281206])
```

```
In [135]: mat
```

```
Out[135]: matrix([[ 0.99992166, -0.01251662],
                  [ 0.01251662,  0.99992166]])
```

```
In [136]: vector * mat
```

```
Out[136]: matrix([[200.01914411, -2.22253163]])
```

```
In [137]: R.I
```

```
Out[137]: matrix([[200.    ,    2.5   ],
                  [   2.5   ,    0.3125]])
```

```
In [ ]:
```